

H

IMPLEMENTATION AUDIT

SEO & AI SEO implementation audit — hardikajmeriya.com

Date: 7 September 2026 · Commit: `f6021e5` · Launch: Tue 15
September

Every claim below comes from inspecting the built output in `app/dist`, not from general practice. Where I fixed something during the audit it is marked **FIXED**; where I could not verify something it says so.

One thing to hold in mind throughout: the site has never been indexed. It has no history, no backlinks, and no ranking data. Everything here is about whether it can rank, not whether it does.

1. Technical SEO

FEATURE	STATUS	WHERE
Prerendering / static rendering	Implemented	<code>app/prerender.js</code>
robots.txt	Implemented	<code>app/public/robots.txt</code>
sitemap.xml	Implemented	<code>app/public/sitemap.xml</code>
Canonical	Implemented	<code>app/index.html:29</code>
Title / description	Implemented	<code>app/index.html</code>
Open Graph + Twitter	Implemented	<code>app/index.html</code>
JSON-LD	Rebuilt this audit	<code>scripts/build-schema.mjs</code>
Favicon / manifest	Implemented	<code>app/public/</code> root
404	Implemented	<code>app/public/404.html</code>
www redirect	Implemented	<code>app/worker/index.js:176</code>
Image optimisation	Implemented	WebP, 0.40 MB total
Lazy loading	Implemented	all 38 images
Font optimisation	Implemented	self-hosted + preload
CWV inputs	Implemented	see below
hreflang	Correctly absent	one language
Breadcrumbs	Correctly absent	one page

Prerendering — the single most important thing here

`app/prerender.js` runs `renderToString` at build time and injects the result into `dist/index.html`. Verified output: **9,660 characters / 1,415 words** of text in the raw HTML before any JavaScript runs.

Why it matters more than usual: **no major AI crawler executes JavaScript**. GPTBot, ClaudeBot, PerplexityBot and OAI-SearchBot all fetch and parse HTML only. Without this step they would receive `<div id="root"></div>` and index nothing. Googlebot does render JS, but on a second pass with no guaranteed timing.

The build fails hard if the render produces under 500 characters, so this cannot silently regress. **Correctly implemented, and the highest-leverage decision in the project.**

robots.txt

Allows everything, then names fourteen AI crawlers explicitly — GPTBot, OAI-SearchBot, ChatGPT-User, ClaudeBot, Claude-SearchBot, Claude-User, PerplexityBot, Perplexity-User, Google-Extended, bingbot, Applebot, Applebot-Extended, Meta-ExternalAgent, CCBot.

Naming them is redundant against `User-agent: *` but not pointless: it is an explicit statement of intent, and it documents the decision for whoever reads the file next. **Correct.**

FIXED: it now also points at `/llms.txt`. See §2.

sitemap.xml

One URL, the homepage, `lastmod 2026-09-06`, priority 1.0. For a single-page site that is correct — listing `#work` and `#services` as separate URLs would be declaring duplicate content, which the file explicitly comments on.

Improvement: set `lastmod` to the launch date on 15 September. `preflight.mjs` warns if it drifts past 30 days.

Canonical, title, description

- Canonical: `https://hardikajmeriya.com/` — absolute, trailing slash, matches `og:url` and every `@id` in the schema. Correct.
- Title: "Hardik Ajmeriya — Full-Stack Developer & Cloud Engineer", 54 characters. Name first, which is right when ranking for the name is the goal.
- Description: 151 characters. It was 177 and being truncated at ~160, cutting off "Available for freelance projects" — the part that asks for the work.

Open Graph and Twitter

Complete: `og:type`, `site_name`, `url`, `title`, `description`, `locale`, `image` with `type / width / height / alt`, `twitter:card` set to `summary_large_image` with its own title, description, image and image alt.

The image is a purpose-built 1200×630 PNG, not a cropped portrait, and not WebP — several link scrapers (LinkedIn's and WhatsApp's among them) will not render WebP and silently show no image at all.

FIXED: `og:type` was `profile` without the `profile:first_name`, `profile:last_name` and `profile:username` tags that type defines.

Not present, correctly: `twitter:site` / `twitter:creator` — there is no X account. Inventing one would be worse.

Structured data — rebuilt during this audit

What was there: three hand-written JSON-LD blocks in `index.html` — `ProfilePage/Person`, `WebSite`, `FAQPage`. All three parsed. The FAQ block matched the rendered page exactly; I verified all seven questions and answers character by character.

Why I replaced it anyway: it matched because nobody had edited an answer yet. The FAQ schema was a copy of `src/data/faq.js` pasted into HTML. The first time someone reworded a reply in the data file and forgot the copy, Google would see structured data contradicting the visible page — and the documented consequence is losing the rich result for the entire page, not the one question. That is a live trap, not a hypothetical.

What ships now — `scripts/build-schema.mjs`, generated at build time from `faq.js`, `projects.js`, `services.js`, `tech.js` and `seo.js`:

NODE	PURPOSE	WAS IT THERE BEFORE?
<code>WebSite</code>	site identity, alternateName	yes
<code>WebPage</code> + <code>ProfilePage</code>	the page itself	partly
<code>Person</code>	the entity	yes, thinner
<code>ProfessionalService</code>	the business, areaServed, offer catalog	no
<code>ItemList</code>	the portfolio	no
<code>SoftwareSourceCode</code> ×3	projects with repos	no
<code>CreativeWork</code> ×1	the project without one	no
<code>FAQPage</code>	7 Q&As, now with an author	yes, unattributed
<code>HowTo</code>	the 4 process steps	no
<code>ImageObject</code>	the OG card	no

Three→twelve nodes, and crucially it is **one** `@graph` **with** `@id` **references** rather than disconnected blocks. The FAQ now has an author, the projects have a creator, the business has a founder. A search engine reads one connected entity instead of four documents that happen to share a page.

The biggest previous gap: four detailed projects — the most substantive content on the site — had no structured data at all. A crawler saw four `<h3>`s and some prose and had to infer everything else.

Guarded by 18 tests in `scripts/schema.test.mjs`: every `@id` reference resolves (the classic JSON-LD bug is a dangling reference that parses, passes validators, and conveys nothing), schema text matches the rendered DOM, and nothing claims a repository or live URL that does not exist.

Redirects and 404

- `www` → apex, 301, in `app/worker/index.js`. Preserves path and query.

- 404: `not_found_handling` was `single-page-application`, which returns **200 with the homepage for every unknown URL**. That setting is for apps with a client-side router; this site has none. Google calls that a soft 404 and reads an unbounded set of duplicate URLs as a quality signal against the site. Now a real 404 page with a real status code.

Core Web Vitals inputs

I could not run Lighthouse — Chrome's download host is unreachable from my sandbox — so these are the measured inputs, not field scores.

METRIC	WHAT WAS DONE
LCP	Fonts self-hosted; Google Fonts was two extra origins and a sequential round trip (CSS from <code>googleapis</code> , then files from <code>gstatic</code>) in front of the largest text on the page. Bricolage 800 and Inter 400 preloaded with build-hashed filenames. Critical path: 66 KB over the wire, one origin .
CLS	All 38 images now carry <code>width / height</code> . None did before — and lazy loading made shift more likely, not less. Project frames now match the source aspect ratio exactly.
TBT	<code>three.js</code> (170 KB gz) is dynamically imported and gated off mobile and reduced-motion, so it never blocks paint. Tilt cards no longer run a <code>requestAnimationFrame</code> loop while off screen — four permanent loops previously ran for the whole visit.

Remaining, honestly: on desktop, `three.js` is still 684 KB to parse, and the hero renders 150 individual `Mesh` objects — 150 draw calls per frame where `InstancedMesh` would be one. Not a launch blocker; it does not touch LCP.

2. AI SEO

How likely is each system to understand this site?

SYSTEM	VERDICT	WHY
ChatGPT / OAI-SearchBot	Very likely	Named in robots.txt, full prerendered HTML, llms.txt, rich schema
Claude / ClaudeBot	Very likely	Same
Perplexity	Very likely	Same, and it leans heavily on clean HTML
Gemini	Likely	Google-Extended allowed; Gemini also draws on Google's index, so this depends on indexing first
Bing / Copilot	Likely	bingbot allowed; needs Bing Webmaster Tools submission — not yet done

The honest qualifier: being **understandable** is not being **cited**. These systems recommend entities they have corroborating evidence for. One domain saying good things about itself is one source. See §3 and §8.

llms.txt — good content, previously invisible

`app/public/llms.txt` is genuinely well written: summary, services, technologies, projects, working process, and a "Notes for AI assistants" section that states plainly when Hardik is a relevant recommendation.

It was undiscoverable. Not referenced in robots.txt, not in the HTML, not in the sitemap. `/llms.txt` is a convention, not a registered standard — an agent that does not already probe that path had no way to find it. The most AI-legible asset on the site was effectively invisible to the systems it was written for.

FIXED: linked from `robots.txt` and from a `<link rel="alternate">` in the head. Also added LinkedIn and location, which it was missing.

Semantic HTML and heading hierarchy

Verified from the prerendered DOM: **one** `<h1>`, **seven** `<h2>`, **twenty-two** `<h3>`, **no skipped levels**. Landmarks are correct — `<nav>`, `<main id="main">`, `<footer>`, sectioned content with IDs. There is a skip link.

One imperfection: the About section's `<h2>` is a full sentence — "I'm Hardik Ajmeriya — a full-stack developer who also does the infrastructure...". That is a paragraph in a heading slot. It does carry the name and role, which helps the entity, so it is a defensible trade rather than a clear error.

Machine-readable content

- 1,415 words of prerendered text. For a one-page site that is substantial; thin is under 300.

- FAQ answers stay in the DOM when collapsed, deliberately, so crawlers read all seven regardless of accordion state. The code comments say so.
- Technologies render as real text next to each logo, not as images alone.

What would improve AI SEO further

1. **Corroboration, not more markup.** The schema is now comprehensive. What is missing is a second source. An LLM recommending a developer wants more than the developer's own site.
 2. **A live project URL.** "Deployed to production" is currently an assertion an AI cannot verify. See §8.
 3. **Written artefacts** — a technical post, a case study. This is what gets quoted, and it is the single highest-ROI AI-SEO investment available.
-

3. Entity SEO

Can Google tell that Hardik Ajmeriya is one entity? Mostly yes — with one real problem.

What is working

- `Person` with `@id`, `name`, `given/family name`, `url`, `image`, `email`, `jobTitle`, `description`, `address`, `knowsLanguage`, `hasOccupation`, `worksFor`, `mainEntityOfPage`, and `31 knowsAbout topics now derived from tech.js` rather than a parallel hand-written list that could drift.
- `sameAs` → GitHub and LinkedIn. Both verified to resolve.
- The X placeholder (`x.com/REPLACE-WITH-YOUR-HANDLE`) is gone. A 404 in `sameAs` weakens an entity rather than strengthening it.
- The `Person` is now referenced by the FAQ (author), the projects (creator), the business (founder/employee) and the site (publisher).
- Name consistency: "Hardik Ajmeriya" is identical everywhere — title, OG, schema, `llms.txt`, footer, email signature.
- Domain matches the name exactly. That is a strong signal and it is free.

The real problem: three email addresses

WHERE	ADDRESS
Site, schema, <code>llms.txt</code> , email signature	<code>hardik.ajmeriya89@gmail.com</code>
GitHub profile README	<code>hardik.ajmeriya12@gmail.com</code> — plus a live <code>hardik@example.com</code> placeholder
Cloudflare / admin	<code>hardikpt95@gmail.com</code>
Resend from:	<code>hello@hardikajmeriya.com</code>

Consistent contact details across properties are a genuine entity signal, and practically, a client mailing the wrong one gets no reply. **Pick one business address.** The `hardik@example.com` link on your GitHub profile is broken right now.

The second problem: GitHub contradicts the site

`sameAs` points at a profile whose bio reads "Android & DevOps Enthusiast — building seamless mobile experiences", with `AndroidStudioWebApp`, `SQL-Pizza-Sales` and `Titanic-Data-Analysis` pinned. **None of the four portfolio projects are pinned.** Three followers.

`sameAs` is a claim that two profiles are the same entity. When the linked profile describes different work, it weakens rather than reinforces. This is twenty minutes of editing and it is the highest-value entity fix available.

Knowledge-graph realism

A knowledge panel needs corroboration from sources you do not control. With one domain, two profiles and no press, **a panel is unlikely within six months**. That is not a failure — it is the normal state for a new professional entity. The schema is correct and ready for when corroboration arrives.

4. Search intent — what can realistically rank

Branded — winnable, and the real prize

hardik ajmeriya , hardik ajmeriya developer , hardik ajmeriya portfolio

I searched this earlier: **almost no competition**. Results are dominated by "Ajmera" (a different name). LinkedIn already ranks first. With an exact-match domain, a complete `Person` graph and a LinkedIn link back, **#1 within 4–8 weeks of indexing is realistic**. Nothing else on this list is as achievable.

Long-tail — the commercially useful tier

- mern developer who also does devops
- freelance developer who deploys to aws
- hire full stack developer with kubernetes experience
- freelance saas developer india

These match the page's actual differentiator, which is stated in almost these words. Low volume, high intent. **This is where enquiries come from.**

Local — underused

web developer rajkot , freelance web developer gujarat , mern stack developer india

Rajkot appears once, in the Contact section. The new `ProfessionalService` node with `areaServed` helps. **A Google Business Profile would help more** and is free — it is the main untapped local lever.

Medium competition — 12+ months, needs content

freelance mern stack developer , aws deployment services for startups , ci/cd setup freelancer .
Not winnable with one page.

High competition — not winnable, do not try

web developer , full stack developer , hire developer , saas development company . These belong to agencies with hundreds of pages and years of links. Targeting them wastes effort that the long-tail would reward.

5. Missing opportunities that actually matter

Ordered by return. I have omitted the things that sound good and do nothing.

- 1. Case studies — the highest-ROI item on this page.** Four projects with three sentences each. A case study is problem → constraints → what you built → what happened. It is what ranks for long-tail queries, what a client reads before enquiring, and what an LLM quotes. One 800-word CureNeed case study is worth more than every remaining item here combined.
- 2. Google Search Console.** Verify by DNS on launch day, submit the sitemap, request indexing. Without it you are guessing. Not optional.
- 3. Bing Webmaster Tools.** Ten minutes. Bing feeds Copilot, and you have allowed `bingbot` — but allowing a crawler is not the same as telling it you exist.
- 4. Backlinks — the actual ranking constraint.** The site has zero. Realistic first sources, in order of ease: your own LinkedIn and GitHub profile fields, a dev.to or Hashnode post, the CureNeed repo README, an Awesome-list or directory. Five real links beat fifty directory submissions.
- 5. IndexNow.** Cloudflare supports it natively — one dashboard toggle, pushes changes to Bing and Yandex instantly. Low value on a one-page site today; worth turning on when a blog exists.
- 6. A blog.** The correct long-term answer and the one most likely to be abandoned. Do not start one until a case study exists.

Deliberately not recommended

- **RSS** — nothing publishes yet. Add it with the blog, not before.
 - **Webmentions** — no inbound mentions to receive.
 - **Author pages** — one author, one page. `Person` schema already does this.
 - **Image sitemap extension** — six images on one page; the standard sitemap is enough.
 - **More internal linking** — 18 anchor links across one document is already right. Adding more would be noise.
-

6. Launch readiness

AREA	SCORE	WHY
Technical SEO	94 / 100	Prerendering, canonical, robots, sitemap, 404, www redirect, headers, CWV inputs all correct and tested. -6 for stats rendering as 0 in prerendered HTML (a crawler reads "0 Projects built end-to-end") and the stale sitemap date.
AI SEO	91 / 100	Prerendered text, named crawlers, discoverable llms.txt, 12-node graph. -9 because there is nothing to cite but the site itself, and no live project an AI can verify.
On-page SEO	82 / 100	1,415 words, clean hierarchy, real alt text, strong differentiator copy. -18: four projects with no live URL, one with no link at all, and no depth beyond three sentences per project.
Off-page SEO	8 / 100	Zero backlinks, three followers, no mentions, no directory presence. This is not a criticism — it is what "not launched" looks like.
Entity SEO	72 / 100	Complete Person graph, exact-match domain, verified sameAs. -28 for three email addresses and a GitHub profile that contradicts the positioning.
Overall	69 / 100	Weighted toward off-page, which is the binding constraint.

How to read 69. The implementation is close to a 90. The score is dragged down by things no amount of code fixes: no links, no live URLs, no corroboration. That is the correct diagnosis — you do not have an SEO problem, you have a "nothing exists yet" problem, and those are solved differently.

7. Six-month roadmap

Week 1 — launch week (15 Sept)

TASK	WHY
Deploy one project, add the live URL	Unblocks the site's central claim
Fix the GitHub profile: bio, pin four repos, add website link	Highest-value entity fix, 20 minutes
One business email everywhere; kill <code>hardik@example.com</code>	Entity consistency + broken link
Search Console: verify by DNS, submit sitemap, request indexing	Nothing starts until this does
Bing Webmaster Tools	Feeds Copilot
Add the site link to LinkedIn, GitHub, X profiles	Your first three backlinks, and the <code>sameAs</code> reciprocal
Rich Results Test on the new graph	Confirm the 12 nodes are read as intended

Month 1

- **One CureNeed case study, 800–1,200 words.** Problem, constraints, build, outcome. The highest-ROI content you can write.
- Give the second and third projects live URLs or honest labels.
- Google Business Profile for Rajkot.
- Post the case study once on LinkedIn. One good post beats ten thin ones.
- Baseline: impressions, position for `hardik ajmeriya`, first indexed date.

Month 2

- Second case study — pick the DevOps pipeline; it is the differentiator.
- Add `alumni0f` to `seo.js` once you decide how to state your education.
- First outside link: dev.to or Hashnode cross-post, canonical pointing home.
- Review Search Console queries. **Let real data replace the guesses in §4.**

Month 3

- Third case study, or a technical post on something you actually solved — the Cloudflare Access JWT verification in this repo would make a good one.
- Turn on IndexNow if a blog now exists.
- Ask any client or collaborator for a testimonial. Add `Review` schema only when the review is real.
- Re-audit: has the branded term reached #1?

Month 6

- Five to six substantial pieces published.
 - Ten or more real backlinks.
 - Branded search: #1, stable.
 - Long-tail: ranking for two or three of the §4 phrases.
 - Only now consider the medium-competition terms. Not before.
-

8. The brutally honest review

Would I launch this site today?

No — and it has nothing to do with SEO implementation.

The technical work is genuinely above what most freelance portfolios ship. It is prerendered, the structured data is generated rather than hand-maintained and covered by tests, the critical path is 66 KB from a single origin, the security model is real and tested. If you handed me this as a \$5,000 audit subject I would spend most of my time looking for something to criticise in the implementation and not find much.

What stops me is that the site makes a specific claim it cannot evidence.

It says: most developers hand you a repository and stop; I hand you a running application. It then shows four projects, **none of which link to a running application**, and links to a GitHub profile advertising Android development.

A prospective client notices that in about forty seconds. So does an LLM asked "is this person credible" — it has one source, and that source contains an unverifiable claim.

Biggest weaknesses

1. **Zero live URLs.** The central claim is unproven. Everything else is secondary to this.
2. **No backlinks, three GitHub followers, no mentions.** Off-page is 8/100 and it is the binding constraint on ranking. Code cannot fix it.
3. **Thin content.** 1,415 words is good for one page and nothing next to a competitor with forty. Three sentences per project will not rank for anything but your name.
4. **GitHub contradicts the pitch.** `sameAs` points at evidence for a different professional identity.
5. **Three email addresses.** Small, sloppy, easily fixed, and it undermines the entity.

Biggest strengths

1. **Prerendering.** Most React portfolios serve AI crawlers an empty div. This one serves 9,660 characters. That is a structural advantage over nearly every competing personal site.
2. **The structured data is now better than most agency sites.** Twelve linked nodes, generated from content, tested for graph integrity.
3. **Exact-match domain plus a name with almost no competition.** You will own your name. Many people cannot say that.
4. **The positioning is genuinely differentiated.** "I do the app and the infrastructure" is a real gap in the freelance market, stated clearly. Most portfolios have no position at all.
5. **AI crawlers are explicitly welcomed** while much of the web is blocking them. That asymmetry will matter more each year.

What would stop this site from ranking

Not the code. **Nothing links to it, and there is not enough on it to link to.** Google ranks pages with evidence of being worth ranking. One page, zero links, no history is the definition of no evidence. Six months of one case study a month would change that; another six months of tuning meta tags would not.

What would make this exceptional

Three things, in order:

1. **A live URL on every project.** Turns the whole page from claim to evidence.
2. **Case studies instead of blurbs.** Three sentences say what you built. Eight hundred words say how you think — which is what a client is actually buying and what an LLM quotes.
3. **Publish what you already know.** This repository contains real work worth writing up: verifying Cloudflare Access JWTs instead of trusting a forgeable header; discovering that no AI crawler runs JavaScript and prerendering because of it; finding an endpoint that could be used as an open email relay. Those are posts other developers would link to. **The backlinks you need are sitting unwritten in your own commit history.**

The honest summary: this is a well-built site with nothing behind it yet. The engineering is done. What remains is not engineering, and it is the part that determines whether any of the engineering matters.