

H

IMPLEMENTATION AUDIT

SEO implementation report

Date: 7 September 2026 · **Scope:** everything implementable in code

First, about the screenshot

The image you shared is a **mockup, not a live result**. Tells: the site is not indexed yet, the meta description shown is not the one in your code, and it credits your GitHub with 200+ followers where the real profile has 3. That is fine — it is a clear statement of the goal, and I have used it as the target.

But two things in it **cannot be produced by code**, and no amount of markup will change that. Saying so now so the rest of this report is read correctly:

The knowledge panel. Panels come from Google's Knowledge Graph, which is built from corroborating sources Google trusts — Wikipedia/Wikidata, major press, established databases. Schema.org markup makes an entity legible; it does not make it notable. There is no `<script>` that produces a panel. What markup does is make sure that when corroboration eventually exists, Google already knows which entity to attach it to. That part is now done.

The five sitelinks. Google builds a sitelink block from **indexed URLs**. Your site had exactly one. `#work` and `#about` are fragments, not URLs — they cannot become sitelinks. This is the single most actionable finding in the whole exercise and it is architectural, not a markup gap. See §Sitelinks below.

What changed

1. The CV was invisible — now it is a real URL

Found: `app/public/DevOps_V1.2.pdf` shipped on every deploy and was linked from **nowhere**. No `href` anywhere in the source, absent from the sitemap, absent from `llms.txt`. An asset nothing points at is not crawlable. The filename also said nothing about whose CV it is, which matters because it is the file a recruiter saves to their desktop.

Changed:

- Renamed to `hardik-ajmeriya-resume.pdf`.
- The Worker serves it at a clean `/resume` (200, not a redirect), so the indexable URL is the readable one. `app/worker/index.js`.
- `Content-Disposition: inline` — a recruiter reads it in the browser rather than being forced into a download. Google indexes PDF text either way.
- `X-Robots-Tag: index, follow` set explicitly.
- Linked from the footer, added to `sitemap.xml`, added to `llms.txt`.
- A `DigitalDocument` node in the graph, `about` and `author` both pointing at the Person.

Impact: your CV becomes indexable and attributable. It is also **the site's second URL** — the beginning of any sitelink story.

5 tests, including that `/resume` on the private preview host is still gated behind Access.

2. Person entity completed

Added to `scripts/build-schema.mjs` :

PROPERTY	VALUE	WHY
<code>contactPoint</code>	ContactPoint, business enquiries, worldwide	The structured form of "how do I hire this person"
<code>identifier</code>	PropertyValue → the domain	The canonical disambiguator. There is no ORCID or Wikidata ID to point at yet, and the exact-match domain is the next strongest thing
<code>subjectOf</code>	→ the CV node	Ties the document to the person
<code>worksFor</code>	→ ProfessionalService	Person and practice are one entity, stated explicitly so they are not read as two competitors for the same name
<code>hasOccupation</code>	Occupation + occupationLocation	Job + place, machine-readable
<code>knowsLanguage</code>	en, gu, hi	
<code>image</code>	the About photo	Panels need a photo to attach
<code>mainEntityOfPage</code>	→ the WebPage	

`knowsAbout` (31 topics) is **derived from** `tech.js` , so it cannot drift from the technology grid the page renders.

3. Business entity gains a logo

`ProfessionalService.logo` — the 512×512 brand mark on an opaque background, which is the shape Google expects. Deliberately distinct from `image` , which is the wide 1200×630 OG card. A test asserts the logo is square and is not the same asset as `image` .

4. humans.txt

`app/public/humans.txt` , in the conventional format: who built it, the stack, the rendering approach, and the explicit note that there is no analytics and there are no cookies. Referenced from `robots.txt` alongside `llms.txt` .

Small, but it is a machine-readable authorship claim on the domain.

5. Discovery files are now discoverable

`robots.txt` points at both `/llms.txt` and `/humans.txt` . Both get an explicit `Content-Type: text/plain; charset=utf-8` in `_headers` — without it Cloudflare can serve `.txt` in a way some clients mis-handle.

6. Redundancy removed from the graph

`hasOccupation.skills` was the `knowsAbout` list joined into a string — the same 31 items serialised twice in one document. Removed.

`hasOfferCatalog` repeated an eight-country `areaServed` for each of seven services: ~3.5 KB of JSON saying nothing the parent had not already said. The offers reference the provider by `@id`, so a consumer follows the link.

What I refused to implement, and why

You asked for these. Implementing them would have been markup describing things that do not exist, which your own brief prohibited ("no fake markup, no spam"). Each is enforced by a test that **fails if it is ever added**.

`SearchAction` / `sitelinks search box`

Requires a working search endpoint. **The site has no search**. Google's documentation is explicit that the markup must point at a URL that returns results; declaring one the site cannot serve is invalid markup and is ignored at best. Becomes legitimate the day a search exists.

`BreadcrumbList`

Breadcrumbs describe a position in a hierarchy. **This is one page**. There is no trail. Marking one up would misrepresent the structure to get a visual result — precisely what the guidelines call out.

`award`, `Review`, `aggregateRating`

No awards, no reviews. Adding any of them would be fabricated. `aggregateRating` on a page with no reviews is one of the most commonly manually-actioned structured-data violations.

`alumniOf`

Genuinely wanted, and I **cannot fill it in**. Your GitHub README says "MCA Graduate" in one place and "ECE Graduate" in another. I will not guess at your education. One-line fix in `SEO_TODO.md`.

X/Twitter placeholder in `sameAs`

You wrote "add placeholders if needed". **No**. A URL in `sameAs` that 404s weakens the entity rather than strengthening it — `sameAs` is a claim of identity, and a claim pointing at nothing is evidence against you. I removed an `x.com/REPLACE-WITH-YOUR-HANDLE` placeholder in an earlier pass for the same reason. Add the real handle when the account exists and links back.

Sitelinks: the honest mechanics

Your mockup shows five sitelinks. Here is what stands between you and them.

Google builds sitelinks from indexed URLs. Before today you had one. You now have two (`/` and `/resume`). Anchors are not URLs — `#work` will never become a sitelink row.

The realistic path, and it is the same conclusion the earlier audit reached from a different direction:

<code>/</code>	the homepage	(exists)
<code>/resume</code>	the CV	(added today)
<code>/work/cureneed</code>	a real case study	(needs writing)
<code>/work/devops-pipeline</code>	a real case study	(needs writing)
<code>/work/kubernetes-monitoring</code>	a real case study	(needs writing)

Three case studies would give you five substantial URLs, internal links pointing at them, and enough per-page content to be worth indexing separately. That is the configuration that produces a sitelink block.

Do not create `/about` , `/contact` and `/tech` as separate thin pages to chase this. Splitting 1,415 words across five pages gives five thin pages, and thin pages do not get sitelinks either. Case studies add content; splitting does not.

Performance: the honest cost

You said never reduce performance. This did cost something, so here is the measurement rather than a reassurance:

	BEFORE	AFTER
JSON-LD block	4.2 KB gz (12 nodes)	4.1 KB gz (13 nodes)
Render-blocking, over the wire	21.6 KB gz	23.5 KB gz
Deploy size	2.1 MB	2.1 MB
Deferred JS	unchanged	unchanged

+1.9 KB gzipped on the critical path. That is the whole cost of the richer graph, after removing the redundancy described above. On any connection this is under 20 ms; on a modern one it is unmeasurable. The JSON-LD sits at the end of `<head>` , after the stylesheet link and font preloads, so it never delays discovery of a render-critical resource.

Accessibility is unchanged: one new footer link, inside the existing 44 px touch-target treatment, with visible text and no new ARIA.

Every file touched

FILE	CHANGE
app/worker/index.js	/resume route
app/worker/ admin.test.mjs	+5 tests for /resume
scripts/build- schema.mjs	contactPoint, identifier, subjectOf, CV node, logo; redundancy removed
scripts/ schema.test.mjs	+7 tests, incl. four that fail if fake markup is added
app/src/components/ Footer.jsx	Résumé link
app/public/ humans.txt	new
app/public/hardik- ajmeriya-resume.pdf	renamed from DevOps_V1.2.pdf
app/public/ robots.txt	points at llms.txt + humans.txt
app/public/ sitemap.xml	/resume , lastmod refreshed
app/public/llms.txt	CV link
app/public/_headers	content-type + cache for the txt files and the CV

Expected impact

Correction — FAQ rich results no longer exist

An earlier report of mine said to expect an FAQ rich result. That was wrong. Google **removed FAQ rich results entirely on 7 May 2026**, and had already restricted them to government and health sites in August 2023. **HowTo** rich results were deprecated in the same 2023 announcement.

The markup stays, and this is not a wasted effort:

- ChatGPT, Claude, Perplexity and Gemini parse **FAQPage** and **HowTo** heavily. Question-and-answer pairs with a named author are among the most quotable structures an answer engine can find.
- It costs ~1 KB gzipped and carries no penalty.
- What changed is only the visible Google result, not the machine-readability.

So: keep it for AI SEO, expect nothing from it in the blue links.

Reliable, within weeks of indexing: - The `Person` entity is unambiguous — name, location, occupation, contact, profiles, CV, 31 topics, all linked in one graph. - `/resume` becomes indexable and attributable; it will rank for "Hardik Ajmeriya resume" / "CV", which has near-zero competition. - Rich-result eligibility for FAQ and the project nodes.

Plausible, over months, given the manual work in `SEO_TODO.md`: - #1 for the branded terms. The name has almost no competition. - AI citations: the prerendered HTML, the named crawlers and `llms.txt` remove every technical barrier. What remains is corroboration.

Not affected by anything in this report: - The knowledge panel. Needs third-party notability. - The sitelink block. Needs more indexed URLs. - Ranking generally. Needs links, which need content.