

H

IMPLEMENTATION AUDIT

SEO scorecard

7 September 2026 · after the implementation pass · 119 tests
passing

Scores are against what is achievable for a one-page personal site, not against an agency with 400 pages. Where a score is capped by something code cannot fix, it says so.

Summary

AREA	SCORE	ONE-LINE REASON
Technical SEO	96%	Everything implementable is implemented and tested
Entity SEO	88%	Complete graph; capped by GitHub contradicting the site
AI SEO	93%	No technical barrier left; capped by having one source
Brand SEO	84%	Strong signals, near-zero competition, zero authority yet
Core Web Vitals readiness	90%	Inputs measured and correct; no field data yet
Knowledge Panel readiness	35%	Markup ready, notability absent — and that is normal
AI Search readiness	91%	Crawlable, legible, welcomed; needs corroboration
Google Search readiness	72%	Ready to be indexed; not yet ready to outrank anything but your own name
Overall	81%	The build is done. The evidence is not.

Technical SEO — 96%

What earns it

- Prerendered to static HTML at build time; 9,667 characters of text before any JS runs, with a build that **fails** if the render drops under 500 chars
- Canonical, absolute, matching `og:url` and every `@id` in the graph
- `robots.txt` allows everything and names 14 AI crawlers; points at `/llms.txt` and `/humans.txt`
- `sitemap.xml` with both real URLs
- Real 404 with a real status code — was previously a soft 404 returning 200 with the homepage for every unknown URL
- `www` → apex 301, path and query preserved, 3 tests including two lookalike hosts that must not redirect
- Self-hosted fonts, 2 preloaded with build-hashed filenames
- All 38 images carry `width / height` ; all lazy-loaded and `decoding="async"`
- Full favicon set at the site root, `site.webmanifest` , `manifest-src` in CSP

- HSTS, `X-Frame-Options`, `X-Content-Type-Options`, `Referrer-Policy`, `Permissions-Policy`, `COOP`, and a CSP with no `unsafe-inline` on `script-src`
- Cache: immutable for hashed assets, revalidate for HTML, `no-store` for 404

The missing 4%

- Stats render as `0` in the prerendered HTML — the count-up starts at zero, so a non-JS crawler reads "0 Projects built end-to-end"
 - No field Core Web Vitals data, because nothing is live
 - `sitemap.xml` `lastmod` needs setting to the launch date
-

Entity SEO — 88%

What earns it

`Person` now carries: `@id`, `name`, given/family name, `url`, `image`, `email`, `jobTitle`, `description`, `address`, `knowsLanguage`, `knowsAbout` (31 topics, derived from `tech.js`), `hasOccupation` with `occupationLocation`, `worksFor`, `contactPoint`, `identifier`, `subjectOf` (the CV), `mainEntityOfPage`, `sameAs`.

Referenced by the FAQ as author, the projects as creator, the business as founder and employee, the site as publisher. One connected graph, not islands.

Exact-match domain. Name consistent everywhere. Both `sameAs` URLs verified to resolve — the `x.com/REPLACE-WITH-YOUR-HANDLE` placeholder was removed, because a 404 in `sameAs` is evidence against an entity, not for it.

Why not higher

- **Your GitHub profile describes a different professional.** Bio says "Android & DevOps Enthusiast"; pinned repos are an Android app, a SQL exercise and a Titanic dataset. `sameAs` points Google at a corroborating source that contradicts the claim. **This is the single largest entity deduction and it is 20 minutes of work** — `SEO_TODO.md` item 5.
 - Three published email addresses (item 6).
 - `alumniOf` empty — I will not invent your education (item 7).
 - Two profiles is a thin `sameAs`; three or four verified ones would be better.
-

AI SEO — 93%

What earns it

- **Prerendering.** No major AI crawler executes JavaScript. GPTBot, ClaudeBot, PerplexityBot and OAI-SearchBot fetch HTML and parse it. Without this they would receive `<div id="root"></div>`. Most React portfolios do exactly that.

- 14 AI crawlers named and allowed while much of the web blocks them
- `llms.txt` — comprehensive, and now actually **discoverable**: it was referenced from nothing, so only a crawler that guessed the path found it
- 13 linked schema nodes; the FAQ has a named author with a job title and a location, which is the expertise signal an answer engine looks for
- 1,415 words of clean semantic HTML: one `h1`, no skipped heading levels, proper landmarks, real text beside every technology logo
- Collapsed FAQ answers stay in the DOM so all seven are readable regardless of accordion state

Why not higher

Nothing technical is left. The cap is that **an AI recommending a developer wants more than that developer's own website**. One source is one source. The 7% is corroboration, not markup.

SYSTEM	WILL IT UNDERSTAND THE SITE?
ChatGPT / OAI-SearchBot	Very likely
Claude / ClaudeBot	Very likely
Perplexity	Very likely
Gemini	Likely — depends on Google indexing first
Bing Copilot	Likely — needs Bing Webmaster submission , TODO item 2

Brand SEO — 84%

What earns it

- Title leads with the name; `og:site_name`, `WebSite.name` and `alternateName` all reinforce it
- Domain is the name, exactly
- H1, About heading, footer, OG card, CV filename and email signature all say "Hardik Ajmeriya" without stuffing
- Location, occupation and stack are stated in prose and in schema
- **Almost no competition**. Search results for the name are dominated by "Ajmera", a different surname. Your LinkedIn already ranks first.

Why not higher

Brand search is winnable but not yet won: no indexing history, no authority, and the second-strongest property in your name (GitHub) currently argues for a different professional identity.

Core Web Vitals readiness — 90%

Not a Lighthouse score. I could not run Lighthouse — Chrome's download host is unreachable from my sandbox — and I will not invent one. These are the measured inputs.

STATE	
LCP	Fonts self-hosted (was 2 extra origins + a sequential round trip), 2 preloaded. Critical path 23.5 KB gzipped from one origin
CLS	All 38 images have intrinsic dimensions; project frames match source aspect ratio
TBT	three.js dynamically imported, gated off mobile and reduced-motion; tilt cards no longer run rAF while off screen
Deploy	2.1 MB, down from 31 MB

The missing 10%: on desktop three.js is still 684 KB to parse, and the hero renders 150 individual `Mesh` objects where `InstancedMesh` would be one draw call. Neither touches LCP. And there is no field data until real users arrive.

Knowledge Panel readiness — 35%

Deliberately the lowest score, and it is not a criticism.

A panel is generated when Google's Knowledge Graph holds an entity it trusts. Trust comes from sources you do not control: Wikipedia/Wikidata, press, established databases, conference listings. **Schema.org markup makes an entity legible; it does not make it notable.** There is no markup that produces a panel.

What is ready (the 35%): a complete, internally consistent `Person` with a photo, location, occupation, contact point, an identifier, a CV, two verified profiles and 31 topics — all under a domain that is the person's exact name. When corroboration appears, Google will already know what to attach it to.

What is missing (the 65%): any third-party source. Realistically **12+ months**, and only if you publish, speak, or get written about. The mockup's panel is a long-term aspiration, not a launch target.

AI Search readiness — 91%

Highest score after Technical, and correctly so — this is where the site is unusually strong for its size. Prerendered content, named crawlers, a discoverable `llms.txt`, an authored FAQ, and project nodes typed as `SoftwareSourceCode` with real repository URLs.

The gap is that **"deployed to production" is currently unverifiable**. Four projects, no live URL. An AI asked whether this developer ships to production finds a claim and no evidence.

Google Search readiness — 72%

Ready to be **indexed**: yes, completely. Nothing blocks a crawler.

Ready to **rank**: only for your own name.

- Zero backlinks
- Two indexed URLs
- 1,415 words total
- No indexing history

That is not a flaw in the build — it is what "never launched" looks like. Every item in `SEO_TODO.md` moves this number; nothing in the codebase does.

Would I launch this to compete for "Hardik Ajmeriya"?

Yes — for the branded terms specifically. The name has almost no competition, you own the exact-match domain, the entity graph is complete and tested, and LinkedIn already ranks for it. **#1 within 4–8 weeks of indexing is a realistic expectation**, provided Search Console submission and the GitHub fix happen (TODO items 1 and 5).

For everything else, not yet — and the reason is unchanged from the earlier audits: four projects, no live URL. It is the one thing that turns the site's central claim from an assertion into evidence, and no schema can substitute for it.

The three highest-return actions, all outside the codebase:

1. Deploy one project and add the live URL
2. Fix the GitHub profile — 20 minutes, biggest single entity gain available
3. Search Console + Bing on launch day